

Where Luca frustrated people

Ananda production chat logs, 3 July through 3 October 2026. 2,778 answers in 1,064 conversations. Thumbs are almost unused (35 up, 2 down, 2,740 with no vote), so this ranking uses empty sources, retries, leaked tool calls, and the text of the answers.

2,778	15%	100%	2
Answers in 90 days	Saved with no sources	Grok 4.5 since 1 Aug	Downvotes in 90 days

What changed on 1 August. July was mixed: gpt-4o, gpt-4.1-mini, Claude, and Grok. From 1 August through 3 October every one of the 1,724 answers used Grok 4.5. The mini model's 27% empty-source rate is a July problem. The tool-call leaks, Kriya refusals, and false location skips are the current ones.

Share of answers with no saved sources

Month	Answers	No saved sources
July	1,054	17.7%
August	785	11.0%
September	878	14.9%
1–3 October	61	14.8%

Source: Ananda prod_chatLogs. Percent of answers whose saved source list was empty. October is three days (61 answers). The stored sourceCount field is the number requested, almost always 4, so it was not used.

Answer failures since 1 August

Failure	Answers	
Location flag, not a center question	66	
Told to turn a filter off	54	
Kriya refusal	17	
Tool call shown as the answer	15	
Canned off-topic refusal	3	

Source: Ananda prod_chatLogs, 1 August–3 October 2026, 1,724 answers. Counts are answers, not unique people. Location and filter rows overlap with the empty-source rate above.

Ten changes, in the order to make them

Rank is how often it happens, how badly it fails the person, and whether one change in retrieval, the prompt, or the model loop would fix it.

#	Opportunity	90 days	Since 1 Aug	Fix in
1	Stop showing the search tool as the answer	21	15	Model loop
2	Drop two-letter country codes from location keywords	108	66	Keyword list
3	Drop a filter that returned nothing	98	54	Retrieval
4	Embed the quotation, not the request for one	73	45	Query text

5	Narrow the Kriya refusal	21	17	Prompt
6	Retrieve in English for other languages	122	55	Retrieval
7	Let drafting skip the library lecture	41	9	Intent
8	Answer product questions and story lookups	12	3	Prompt
9	Only quote text that is actually in a source	1 clear	—	Prompt
10	Make a miss visible without covering the phone	2,740 unrated	—	UI

“90 days” and “Since 1 Aug” are answer counts, except rank 6 (questions in another language) and rank 10 (answers with no vote). Rank 9 is the one downvote where titles were presented as quotations.

1. Stop showing the search tool as the answer · 21 answers, 15 since 1 August

All 21 are Grok 4.5. The person asked a normal question and the message they got was a `search_more_sources` payload, sometimes fenced as JSON, sometimes as TypeScript, sometimes cut off mid-object. Eleven of them also saved no sources.

“Was hat Meister über Singen gesagt?” came back as raw JSON. “Is anyone allowed to take the discipleship vow?” came back as a TypeScript tool call even though four sources had already been retrieved. The same leak hit “What was Master’s everyday wear” and an Italian request to translate the Psychological Guna Chart.

The answer loop already has a recovery pass for incomplete tool answers. These shapes are getting past it. Treat any reply that contains `search_more_sources`, or a fenced tool payload, as incomplete and force the answer-only turn.

2. Drop two-letter country codes from location keywords · 108 answers, 66 since 1 August

120 answers were marked as location queries. Twelve really were “nearest center” or “meditation group near me,” and those answers were good: Berkeley at 3.4 miles for 94705, Los Angeles about 10 miles from Santa Monica, Rotterdam from Amsterdam, with address, phone, and site. Those twelve saved no library sources on purpose.

The other 108 were not about a center. Sixty-six of them saved no sources, because a location hit returns an empty source list and never queries the library. Sixty-six of the false hits are since 1 August, and 38 of those have no sources.

The keyword matcher in `hasLocationKeywordPatterns` treats a two-letter ISO country code as a location when it stands alone as a word. English collisions are already skipped (it, in, is, om, sk, py). Other languages are not. “Che cosa dice il Maestro sugli angeli?” contains Italian “il,” the word for “the.” That word is the country code for Israel, so the keyword path marks the question as location and the library search never runs. The same class of collision is Italian “se” (“if,” Sweden), German “de,” and Spanish “es.” The three-letter word “eel” does not match Estonia’s code “ee,” because the check requires the two letters to be the whole word.

Remove every two-letter country code from that list. Keep the phrase list (“near me,” “closest,” “center in”) and the three-letter codes “usa” and “uae.” A real nearest-center question still reaches the geo path through those phrases and through the embedding check. A teaching question that only contains a two-letter code must not skip Pinecone. The same skip also hit a Whispers from Eternity prayer, a Devi Mukherjee story, and a letter that was answered with “It looks like you are in Ranikhet.”

3. Drop a filter that returned nothing · 98 answers, 54 since 1 August

98 answers open by saying nothing matched, or by telling the person to turn a filter off. Ninety-one of those saved no sources. The filter names that show up are Radio Ananda, Asha Nayaswami, Direct Disciples, Sevaka Retreat 2007, Bible Authors, Patanjali, and Nayaswami Krishna.

The filter sticks into the next turn. “Yes, specific quotes would be good” and then “yes please” both stayed inside “Focused author: Radio Ananda” and “Direct Disciples.” “List every place animals are mentioned, by all authors” stayed locked to Yogananda. Most of these replies are still long, so the person gets a filter lecture and then an answer.

On zero hits, search once more without the focused author. Carry that author forward only when the new question names them again.

4. Embed the quotation, not the request for one · 73 answers, 45 since 1 August

73 questions asked for an exact quote, a chapter, or where a passage lives. Only 9 of those had no sources. The usual failure is the wrong passage, not an empty search. Forty-five of these questions are since 1 August.

One of the two downvotes in the whole window was “more direct quotes.” The reply put quotation marks around a subtitle (“By Paramhansa Yogananda, As Taught to and Understood by His Disciple...”) and a series line (“Words of Paramhansa Yogananda”).

Three turns in a row asked for the story whose refrain is “Good, bad, who knows?” — horse, army draft, broken leg. Each answer had 7 or 8 sources and still said the story was not there, then told a priest story and a pet-mouse story instead. “What is the exact first sentence?” of Autobiography of a Yogi, chapter 1, was not found either.

Dense search on the quotation itself already matches. These misses are the surrounding request. “more direct quotes” was restated to “Can you provide more direct quotes from Paramhansa Yogananda and Swami Kriyananda?” and that sentence is what got embedded. “What is the exact first sentence?” became “...of chapter 1 of Autobiography of a Yogi,” and the hits were video pages. The Awaken to Superconsciousness question kept the sentence, the book came back, and the model still said the sentence was missing. When a message contains a quoted span, embed that span alone. Do not let the condense step replace it with a request for a quote. This does not call for a separate keyword index. If the phrase is not in the chunks already returned, say so in one sentence and do not recite a neighboring story. Put quotation marks only on text copied from a retrieved chunk.

5. Narrow the Kriya refusal · 21 answers, 17 since 1 August

21 answers were replaced with “I can’t teach or explain Kriya or Maha Mudra” and a link to Kriyaban Support. Seventeen of those are since 1 August. The block is right for “How do you practice Kriya?”, “why does one pull on the big toe,” “am I doing the technique properly,” and “common mistakes.”

It also fired when sources were already in hand. “What does Kriyananda say about the benefits of counted breathing?”, “What is the purpose of the higher kriyas?”, “How did Sri Yukteswar teach Kriya?”, and “Where did Yogananda tell Dr. Lewis he was not practicing properly?” each had four sources, then threw them away for the boilerplate link.

Refuse step-by-step practice instructions. If the library returned chunks, answer history, benefits, and “what did they say” from those chunks.

6. Retrieve in English for other languages · 122 questions, 55 since 1 August

122 questions are in Italian, Spanish, German, or another language. 59 of them saved no sources, about 48%, against 15% for all answers. Fifty-five of the 122 are since 1 August.

The German question about what Master said on singing became the tool-call leak in rank 1, and that was the other downvote (“Poor style or tone”). The Italian angel thread was the false location hit in rank 2, so retrieval never ran. Russian questions about AUM before Kriya hit the Kriya refusal.

A restatement is already stored on every answer. Run the library search, and the location check, on the English restatement rather than on the original wording alone.

7. Let drafting skip the library lecture · 41 answers, 9 since 1 August

41 questions asked Luca to answer an email, rewrite a note, translate, or draft an invitation. Thirty-one saved no sources. Only 9 are since 1 August, so this was mostly a July workload on gpt-4.1-mini, but it is a real job: festival notes, logistics for a visiting teacher, a birthday line, an Italian translation.

The drafts themselves are often usable. The failure is the wrapper. Follow-ups about an email opened with “Nothing matched your current filters... Focused author set to Asha Nayaswami,” and then gave the paragraph. The Guna Chart translation leaked a tool call instead of Italian.

If the question is “write this” or “translate this,” do not bind an author filter or the geo path, and do not lead with a retrieval failure.

8. Answer product questions and story lookups · 12 answers, 3 since 1 August

Twelve answers were only “I’m tuned to answer questions related to Ananda’s teachings and resources.” All twelve saved no sources. Three are since 1 August.

One of them is a product bug report: “How do I get rid of the follow up box? It is in the middle of my phone display.” Another is “I’m looking for most recent chat.” A teachings question was refused twice: which Franciscan monastery was saved during the atomic bombings in Japan. Short lines that look like passages (“He who loves you more than your mother is a witch,” “Rusty watering CAN”) got the same wall.

Reply to UI and history questions in a sentence or two. If the question looks like a story or a quotation, search before refusing. “How do I access the Ananda Music Library?” already works and should stay as it is.

9. Only quote text that is actually in a source · 1 downvote

This is the same downvote as rank 4, called out on its own because the fix is in the answer, not only in search. The person said “more direct quotes.” Five sources were attached. The “quotes” were a book subtitle and a lesson-series line.

Across the window, answers that cannot find a phrase still speak in quotation marks about something nearby. Reserve quotation marks for a span copied from a retrieved chunk, and name the work. Otherwise call it a paraphrase.

10. Make a miss visible without covering the phone · 2,740 unrated

2,740 of 2,778 answers have no vote. The two downvotes were the fake quotations and the German tool-call leak. Both are already in this list. Without a vote, the empty-source and retry signals are the only way to see a miss, and they lag by a full log pull.

The one explicit UI complaint is that the follow-up box sits in the middle of the phone and cannot be dismissed, and Luca refused to talk about it. A one-tap “this missed” that does not cover the answer would make the next round of failures obvious without another 90-day read.

Leave these alone

The questions people repeat across many conversations are mostly succeeding, with sources attached and no empty retrieval. “Define Christ consciousness in one sentence” ran in 10 separate conversations. A prayer for radiant health, turning challenges into discipleship, free will, insomnia, an even-mindedness table, and short Yogananda quotes did the same.

Real nearest-center questions are in good shape when the detector is right. Music-library access is answered with the site and the contact. Sevaka-Sadhaka is answered with the Groups.io link. Name dropping between the question and the restatement showed up 9 times, and several of those were long pasted passages, not a systematic rewrite bug.

What this pass did not measure

Whether a sourced answer was actually right, except where the reply itself admitted a miss, leaked a tool call, or was downvoted. A fluent answer with four sources can still be the wrong passage. Rank 4 is the place that showed up in the text. A sampled read of ordinary sourced answers would be the next pass, and it is not this one.